Africa TTS: Fast, CPU-Friendly Speech Synthesis for 92 African Languages from Phonemes

Mich-Seth Owusu
AfriSpeech
michseth@ghananlp.org

Abstract

We present Africa TTS, a text-to-speech system for 92 African languages and English that runs faster than real time on an ordinary CPU. All languages share one input space: text is converted to IPA by rule-based grapheme-to-phoneme tables, so a single model learns sounds rather than orthographies. A 24-layer continuous-latent audio language model, built on Pocket TTS, is trained on the IPA, and a 6-layer student, Africa TTS Turbo (110M parameters), is distilled from it with classifier-free guidance folded into its weights. Africa TTS Turbo generates speech 2.38× faster than real time on four CPU cores and starts streaming within 0.17 s. On held-out sentences in 47 languages, transcribed by Omnilingual ASR, Africa TTS Turbo reaches a mean character error rate (CER) of 9.0% against 10.3% for MMS-TTS (better in 21 of 27 languages) and 5.9% against 10.3% for the 3.3B-parameter Sunbird Orpheus, and comes within 0.4 points of the 813M-parameter OmniVoice (8.8% vs. 8.4%) at about a seventh of its size. Models, code, voices and the benchmark are released.

1 Introduction

Speech synthesis for African languages is held back less by modelling than by scale and deployment. Most of the continent's more than two thousand languages have little or no transcribed speech, and the people who would use a voice interface in those languages often rely on phones and low-cost laptops without a GPU. Massively multilingual systems have widened coverage: MMS-TTS [19] trains one small VITS model [11] per language for more than a thousand languages, and OmniVoice [22] clones voices zero-shot in more than six hundred. But per-language models do not share what languages have in common, and large zero-shot models need accelerators to run fast.

We present **Africa TTS**, a single text-to-speech model for 92 African languages and English, built on Kyutai's Pocket TTS [13, 20], a continuous-latent audio language model designed to run faster than real time on a CPU. Two choices shape the system:

1. **One phonetic input space.** Every language is converted to IPA by rule-based grapheme-to-phoneme tables [1] (espeak-ng [6] for English), so the model learns sounds rather than 92 orthographies, and languages with little data share acoustic knowledge with their neighbours.
2. **Speech for the CPU.** A 6-layer student, **Africa TTS Turbo** (110M parameters including the codec), is distilled from the 24-layer teacher with classifier-free guidance [9] folded into its weights. It generates speech 2.38× faster than real time on four CPU cores and starts streaming after 0.17 s, and ships as int8 ONNX bundles for servers and phones.

We evaluate intelligibility with Omnilingual ASR [17] on held-out sentences in 47 languages, against MMS-TTS, OmniVoice and Sunbird Orpheus, with human recordings of the same sentences as a reference. Our contributions are: (i) a single phoneme-input TTS model covering 92 African languages, trained on open data; (ii) a 110M-parameter distilled model that is fast enough for CPU-only and on-device use, with ONNX exports; and (iii) an open intelligibility benchmark over 47 languages comparing four systems, with all synthesised audio and transcripts.

All models, voices, code, the benchmark and its outputs are released.[1]

2 Related Work

Multilingual TTS for African languages. BibleTTS [15] released studio-quality single-speaker data and VITS baselines for ten sub-Saharan languages. MMS-TTS [19] scaled per-language VITS models to over 1,100 languages from New Testament recordings, using uroman [7] romanisation for non-Latin scripts. OmniVoice [22] is a single 0.6B-class discrete non-autoregressive model trained on 581k hours of open data in over 600 languages, with zero-shot voice cloning. WAXAL [5] recently added ASR and TTS corpora for over twenty African languages, and Sunbird Orpheus [21] fine-tunes a 3B-parameter speech language model on these and other data for a dozen languages of East, West and Southern Africa.

[1] Code: https://github.com/AfriSpeech/africa-tts; models: https://huggingface.co/AfriSpeech/africa-tts-turbo and https://huggingface.co/AfriSpeech/africa-tts; live demo: https://huggingface.co/spaces/AfriSpeech/africa-tts-demo.

Efficient speech language models. Pocket TTS [13] follows the continuous audio language model of Rouard et al. [20]: a causal Transformer predicts, frame by frame at 12.5 Hz, a continuous latent of the Mimi codec [4], sampled by a small consistency-trained MLP head, with voice cloning from a short audio prompt. We keep the architecture, replace the text input with IPA, and distil depth [8].

Evaluation with ASR. Automatic intelligibility is commonly measured as the word or character error rate of an ASR system on synthesised speech. Omnilingual ASR [17] covers over 1,600 languages, which for the first time allows a single ASR to score TTS in most of the languages studied here.

3 Africa TTS

3.1 Data

We train on 897 hours of real speech: JW recordings (383 h), a crowd-sourced female speech corpus in 17 languages (192 h), Ghanaian-accented English (156 h) and a 200-hour slice of HiFiTTS-2 [14] (166 h). Any clip whose transcript contains a numeral is dropped (the model reads spelled-out numbers only), as are clips outside 1–30 s. Word timings for voice-prompt cut points come from ctc-forced-aligner [2] for every language (99.95% of clips aligned). African languages range from 0.04 to 27 hours each (33 have under two hours); English has 323 hours.

3.2 Phonetic front end

Text is converted to IPA word by word: African languages by the C tables of africa-g2p [1], English by espeak-ng 1.51 with normalisation. The same C library runs on device and is byte-identical to the training conversion on every validation sentence. Common input errors are normalised, e.g. Greek and Cyrillic lookalikes of the African letters ɛ and ɔ. A SentencePiece unigram tokenizer [12] with 4,000 pieces is trained on the IPA, sampling languages in proportion to the square root of their hours (2.41 tokens per word).

3.3 Teacher

The 24-layer teacher (336M parameters, of which 20M are the Mimi codec) is warm-started from Kyutai's English Pocket TTS with a fresh text embedding table for the IPA vocabulary, and fine-tuned on one H200 GPU with batch 64, learning rate 2×10^{-4} and 20% text and voice dropout (enabling classifier-free guidance). Training early-stopped at 113k steps; the release uses step 100k, chosen by listening. At inference the teacher runs either *guided* (cfg 2.0: two backbone passes per frame) or *unguided* (one pass).

3.4 Africa TTS Turbo: depth distillation

The 6-layer student is initialised from teacher layers 0, 1, 2, 21, 22 and 23, with the text table, flow head and end-of-speech head copied, and trained to match the teacher's guided output, so guidance costs nothing at inference. It trains for 200k steps (cosine schedule, peak 4×10^{-4}); among checkpoints at 150k, 175k and 200k, 200k was preferred in listening tests. Per frame, Turbo computes 6 layers where the guided teacher computes 48.

3.5 Deployment

Both models export to ONNX and run in sherpa-onnx [10] with int8 weights: 209 MB for Turbo and 437 MB for the teacher, including the front end and default voices. The int8 models match PyTorch within 6.5×10^{-5} before quantisation and the decoder at 116 dB SNR. A CPU-only API server (HTTP, WebSocket streaming) powers a public demo that scales to zero.

4 Evaluation

4.1 Speed

Table 2 reports speed on a dedicated 4-vCPU cloud machine over eight sentences in eight languages. Turbo is the only configuration comfortably faster than real time, and it starts streaming within 0.17 s.

4.2 Intelligibility: setup

Test sentences. For each trained language we take up to 50 sentences from, in order of preference, the Omnilingual ASR corpus [17] (test split, else dev, else train; five languages with tiny splits are topped up from the next split), FLEURS [3] test, WAXAL [5] test, and FLORES+ [16, 18] devtest (text only). The first three come with human recordings. Sentences with digits, under 20 or over 220 characters are removed, and Greek or Cyrillic lookalike letters in FLORES+ are replaced by the intended African letters. None of these sentences is in our training data.

Systems. Africa TTS in its three modes (Turbo, guided, unguided) uses each language's default voice and temperature 0.3. MMS-TTS [19] is run for every language with a released model, with uroman for non-Latin scripts. OmniVoice [22] is run for every language in its list, cloning the same default voice clip as Africa TTS; as OmniVoice also conditions on the prompt's transcript, each prompt is transcribed with the evaluation ASR where no full human transcript exists. Sunbird Orpheus [21], a 3.3B-parameter Llama-based model that generates SNAC codec tokens, is run for the languages in its training mix with one fixed speaker per language and its recommended sampling settings (temperature 0.6, top-p 0.95, repetition penalty 1.1). Parameter counts are measured on the released checkpoints and include each system's codec or vocoder.

		Africa TTS (ours)					
	Real *speech*	**Turbo** *110M*	**Guided** *336M*	**Unguided** *336M*	**MMS-TTS** *36M/lang.*	**OmniVoice** *813M*	**Orpheus** *3.3B*
CER (%)							
Languages MMS-TTS supports (27)		9.0	**8.7**	10.6	10.3		
Languages OmniVoice supports (29)		8.8	8.8	10.4		**8.4**	
Languages Sunbird Orpheus supports (11)		5.9	**5.6**	7.3			10.3
All systems (6)		6.0	5.7	7.2	7.8	**4.3**	9.8
Languages with real recordings (24)	7.5	**7.7**	**7.7**	8.9			
CER without diacritics (%)							
Languages MMS-TTS supports (27)		7.6	**7.3**	9.2	9.4		
Languages OmniVoice supports (29)		7.5	7.5	9.1		**7.0**	
Languages Sunbird Orpheus supports (11)		4.6	**4.4**	6.1			8.8
All systems (6)		4.4	4.1	5.7	6.5	**3.0**	7.7
Languages with real recordings (24)	6.6	6.5	**6.3**	7.6			

Table 1: Intelligibility on languages Africa TTS was trained on: mean per-language character error rate of Omnilingual ASR transcripts (lower is better), averaged over the languages each comparison shares (count in brackets; languages with fewer than 10 test sentences are excluded). Best synthetic system in bold; real speech is the ASR's error on human recordings of the same sentences. Sizes are parameter counts including each system's audio codec or vocoder.

Model	**Runtime**	**×RT ↑**	**TTFA ↓**
Africa TTS Turbo	sherpa-onnx int8	**2.52**	–
Africa TTS Turbo	PyTorch	2.38	0.17 s
Africa TTS unguided	sherpa-onnx int8	0.82	–
Africa TTS unguided	PyTorch	0.90	0.52 s
Africa TTS guided	PyTorch	0.47	0.72 s

Table 2: CPU speed on a dedicated 4-vCPU cloud machine: seconds of audio generated per second of compute (×RT, above 1 is faster than real time) and median time to first audio when streaming (TTFA). PyTorch runs on one thread, sherpa-onnx on four.

Metric. Each output is transcribed by Omnilingual ASR (`omniASR_LLM_7B_v2`, conditioned on the language where it has a code), and we report the character error rate (CER) after Unicode NFC normalisation, lower-casing and removal of punctuation. A second CER also strips diacritics, as several G2P tables and ASR outputs treat tone marks inconsistently. The same ASR on the human recordings gives a reference: synthetic speech can score below it, since read synthetic speech is cleaner than spontaneous or noisy recordings. Means are taken over the languages each comparison shares.

4.3 Results

Table 1 summarises intelligibility. Against MMS-TTS, the strongest widely used per-language baseline, every guided configuration is more intelligible: Africa TTS Turbo averages 9.0% CER against 10.3% and is better in 21 of 27 languages, and the guided teacher in 24 of 27. Sunbird Orpheus, a 3.3B-parameter model, is the least intelligible system on the 11 languages it covers (10.3% against 5.9% for Africa TTS Turbo, which is better in 8 of 11 of them). OmniVoice is the most intelligible system overall: on its 29 shared languages it averages 8.4% against 8.8% for both Africa TTS Turbo and the guided teacher, and it is better in most individual languages, by small margins. Africa TTS Turbo closes most of that gap at about a seventh of OmniVoice's parameters while running faster than real time on a CPU (Table 2). Distillation costs nothing in intelligibility: Africa TTS Turbo matches the guided 24-layer teacher, while the unguided teacher is clearly worse (10.6% on the MMS-TTS languages). On the 24 languages with human recordings, Africa TTS Turbo and the guided teacher (7.7%) are within 0.2 points of the ASR's error on real speech (7.5%). Removing diacritics lowers all CERs but leaves the ranking unchanged.

5 Limitations

The ASR-based metric measures intelligibility only, not naturalness or speaker similarity, and inherits the ASR's own errors; for languages whose test text comes from the Omnilingual ASR corpus's training split, the ASR may have seen those sentences, which favours all systems equally. FLEURS sentences are FLORES sentences read aloud. Three languages (Chitumbuka, Swati, Kimbundu) have no ASR language code and are transcribed unconditioned. Most of our African training speech is read religious text, so conversational and technical registers are less natural, and languages with under about two hours of data are weak. Some G2P tables do not mark tone. Voice cloning can be misused to impersonate people; the models are released for non-commercial use under CC-BY-NC-4.0.

6 Conclusion

Africa TTS shows that one phoneme-input model can serve 92 African languages and that, distilled to 110M parameters,

it runs faster than real time on a CPU without losing intelligibility. It is clearly more intelligible than per-language MMS-TTS models and a much larger Sunbird Orpheus, and close to OmniVoice at a fraction of its size. Future work includes human listening tests with native speakers, more conversational training speech, and tone marking in more G2P tables.

References

[1] AfriSpeech. africa-g2p: grapheme-to-phoneme rules for african languages. https://github.com/AfriSpeech/africa-g2p, 2026.

[2] Mahmoud Ashraf. ctc-forced-aligner. https://github.com/MahmoudAshraf97/ctc-forced-aligner, 2024.

[3] Alexis Conneau, Min Ma, Simran Khanuja, Yu Zhang, Vera Axelrod, Siddharth Dalmia, Jason Riesa, Clara Rivera, and Ankur Bapna. FLEURS: Few-shot learning evaluation of universal representations of speech. In *IEEE Spoken Language Technology Workshop (SLT)*, 2023.

[4] Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou, Edouard Grave, and Neil Zeghidour. Moshi: a speech-text foundation model for real-time dialogue. *arXiv preprint arXiv:2410.00037*, 2024.

[5] Abdoulaye Diack, Perry Nelson, MohamedElfatih MohamedKhair, Subhashini Venugopalan, Emmanuel Asiedu Brempong, Tavonga Siyavora, Bob MacDonald, Uche Okonkwo, Sandy Ritchie, Mandy Jordan, Abhishek Bapna, Daan van Esch, Vusumuzi Dube, Jason Hickey, Ronit Levavi Morad, Yossi Matias, Jeff Dean, Aisha Walcott-Bryant, and Avinatan Hassidim. WAXAL: A large-scale multilingual african language speech corpus. Google Research. https://research.google/pubs/waxal-a-large-scale-multilingual-african-language-speech-corpus/, 2026.

[6] eSpeak NG contributors. eSpeak NG. https://github.com/espeak-ng/espeak-ng, version 1.51, 2022.

[7] Ulf Hermjakob, Jonathan May, and Kevin Knight. Out-of-the-box universal Romanization tool uroman. In *Proceedings of ACL 2018, System Demonstrations*, 2018.

[8] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.

[9] Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. *arXiv preprint arXiv:2207.12598*, 2022.

[10] k2-fsa. sherpa-onnx. https://github.com/k2-fsa/sherpa-onnx, 2026.

[11] Jaehyeon Kim, Jungil Kong, and Juhee Son. Conditional variational autoencoder with adversarial learning for end-to-end text-to-speech. In *International Conference on Machine Learning (ICML)*, 2021.

[12] Taku Kudo and John Richardson. SentencePiece: A simple and language independent subword tokenizer and detokenizer for neural text processing. In *Proceedings of EMNLP: System Demonstrations*, 2018.

[13] Kyutai. Pocket TTS. https://github.com/kyutai-labs/pocket-tts, 2026.

[14] Ryan Langman, Xuesong Yang, Paarth Neekhara, Shehzeen Hussain, Edresson Casanova, Evelina Bakhturina, and Jason Li. HiFiTTS-2: A large-scale high bandwidth speech dataset. In *Interspeech*, 2025.

[15] Josh Meyer, David Ifeoluwa Adelani, Edresson Casanova, et al. BibleTTS: a large, high-fidelity, multilingual, and uniquely african speech corpus. In *Interspeech*, 2022.

[16] NLLB Team, Marta R. Costa-jussà, James Cross, et al. No language left behind: Scaling human-centered machine translation. *arXiv preprint arXiv:2207.04672*, 2022.

[17] Omnilingual ASR team, Gil Keren, Artyom Kozhevnikov, Yen Meng, Christophe Ropers, Matthew Setzler, et al. Omnilingual ASR: Open-source multilingual speech recognition for 1600+ languages. *arXiv preprint arXiv:2511.09690*, 2025.

[18] Open Language Data Initiative. FLORES+. https://huggingface.co/datasets/openlanguagedata/flores_plus, 2024.

[19] Vineel Pratap, Andros Tjandra, Bowen Shi, Paden Tomasello, Arun Babu, et al. Scaling speech technology to 1,000+ languages. *arXiv preprint arXiv:2305.13516*, 2023.

[20] Simon Rouard, Manu Orsini, Axel Roebel, Neil Zeghidour, and Alexandre Défossez. Continuous audio language models. *arXiv preprint arXiv:2509.06926*, 2025.

[21] Sunbird AI. orpheus-3b-tts-multilingual. https://huggingface.co/Sunbird/orpheus-3b-tts-multilingual, 2025.

[22] Han Zhu, Lingxuan Ye, Wei Kang, Zengwei Yao, Liyong Guo, Fangjun Kuang, Zhifeng Han, Weiji Zhuang, Long Lin, and Daniel Povey. OmniVoice: Towards omnilingual zero-shot text-to-speech with diffusion language models. *arXiv preprint arXiv:2604.00688*, 2026.

A Per-language results

Language	Code	Src	n	Real *speech*	**Africa TTS** (ours) **Turbo** *110M*	**Guided** *336M*	**Unguided** *336M*	**MMS-TTS** *36M/lang.*	**OmniVoice** *813M*	**Orpheus** *3.3B*
Amharic	amh	FL	50	8.0	5.7	5.6	8.3	10.9	**4.8**	–
Changana (Mozambique)	tso	FR	50	–	11.5	10.7	11.7	**5.0**	–	–
Chitumbuka†	tum	FR	50	–	11.5	**11.4**	14.4	–	–	–
Chokwe	cjk	FR	50	–	14.9	14.6	18.8	–	**12.9**	–
Cibemba	bem	FR	50	–	**5.3**	5.7	11.0	7.3	–	–
Cinyanja	nya	FL	50	8.0	5.0	3.8	4.7	7.1	**2.6**	–
Dagaare	dga	WX	50	22.1	10.4	10.5	**10.2**	10.6	–	–
English	eng	FL	50	2.3	2.4	**2.0**	3.8	3.5	**2.0**	2.3
Ewe	ewe	FR	50	–	6.4	**5.8**	6.6	6.2	–	6.3
Fante	fat	OC	50	13.9	**7.3**	7.7	9.8	–	8.6	–
Fon	fon	FR	50	–	22.5	24.5	24.0	**20.9**	–	–
Frafra	gur	OC	50	5.5	9.8	10.8	11.8	–	**9.1**	–
Hausa	hau	FL	50	4.7	4.2	3.8	4.7	6.2	**3.6**	8.6
Herero	her	OC	14	3.8	4.8	**1.2**	3.2	–	1.5	–
Igbo	ibo	FL	50	8.2	6.9	7.2	8.1	–	**6.6**	17.5
Isoko	iso	OC	5	13.8	29.4	24.5	26.6	–	**14.9**	–
Jula	dyu	FR	50	–	**23.8**	25.1	28.5	29.7	47.9	–
Kabiye	kbp	FR	50	–	11.1	**6.9**	8.8	11.7	–	–
Kabuverdianu	kea	FL	50	2.6	11.9	12.6	**8.3**	–	8.6	–
Kabyle	kab	FR	50	–	14.0	12.6	15.4	14.5	**6.8**	–
Kamba	kam	FL	50	10.3	14.6	13.7	17.4	–	**13.5**	–
Kikuyu	kik	FR	50	–	**7.0**	7.1	7.8	8.3	–	21.3
Kimbundu†	kmb	FR	50	–	**12.5**	13.9	15.4	–	–	–
Kinyarwanda	kin	FR	50	–	3.2	4.2	6.2	6.3	**1.8**	10.4
Kirundi	run	FR	50	–	3.0	**2.8**	3.4	10.2	–	–
Kpelle	xpe	OC	1	21.1	34.1	26.8	30.9	–	**26.0**	–
Lingala	lin	FL	50	3.1	3.8	3.8	9.4	–	**3.4**	3.8
Luganda	lug	FL	50	7.3	8.2	7.0	11.7	8.8	**5.2**	6.6
Macua	vmw	FR	50	–	6.7	**5.9**	10.0	6.0	–	–
Malagasy	mlg	OC	50	3.6	8.8	8.4	14.6	9.6	**5.2**	–
Mauritian Creole	mfe	FR	50	–	7.1	**5.2**	5.9	5.9	–	–
Meru	mer	OC	21	9.2	**14.1**	15.5	16.7	–	25.6	–
Mooré	mos	FR	50	–	15.4	16.1	20.1	**14.1**	–	–
Ndonga	ndo	OC	13	14.0	5.9	4.2	**3.7**	–	5.6	–
Sango	sag	FR	50	–	**10.0**	10.2	11.8	10.6	–	–
Sena	seh	FR	50	–	**6.2**	6.4	7.4	7.3	–	–
Sepedi	nso	FL	50	5.9	4.9	9.4	9.9	–	**4.1**	–
Shona	sna	FL	50	2.9	3.0	2.8	3.2	6.1	**1.9**	–
Swahili	swh	FL	50	1.8	1.2	1.2	1.2	2.6	**0.6**	2.6
Swati†	ssw	FR	50	–	**9.5**	9.9	14.7	–	–	–
Tigrinya	tir	FR	50	–	11.9	11.5	13.7	17.6	**7.7**	–
Twi	twi	FR	50	–	**12.1**	12.7	15.4	13.1	12.5	–
Umbundu	umb	FL	50	9.5	16.6	17.3	15.3	–	**14.8**	–
Wolof	wol	FL	50	10.3	12.2	11.7	13.1	–	**7.2**	–
Xhosa	xho	FL	50	3.4	4.3	3.9	5.3	–	**3.5**	5.9
Yoruba	yor	FL	50	15.9	17.0	15.9	15.8	19.3	**12.8**	28.2
Zulu	zul	FL	50	3.7	**2.2**	3.9	4.4	–	2.4	–

Table 3: CER (%) per trained language. Src: OC Omnilingual ASR corpus, FL FLEURS, WX WAXAL, FR FLORES+ (text only, no real recordings). n: sentences. †: no Omnilingual ASR language code; transcribed without one. Best synthetic system in bold.